

# DSCD 614 - Reinforcement Learning

## Programming Assignment 1: Frozen Lake from First Principles Using Q-Learning

Name: Nyamekye Emmanuel Barima Student ID: 22425704

### 1. Introduction

This project implements a complete reinforcement learning solution for the Frozen Lake problem from first principles. Reinforcement Learning is a learning approach in which an agent improves its behaviour by interacting with an environment, receiving rewards and updating its future decisions. In Frozen Lake, the agent starts at S and must reach G while avoiding holes H. The implementation uses Python with NumPy and Matplotlib only; no reinforcement learning framework is used.

### 2. Environment Design

The environment is an 8 by 8 grid-world. Each cell is either Start, Frozen, Hole or Goal. The required environment methods `reset()`, `step(action)`, `render()`, `get_state()` and `is_terminal()` are implemented in a custom `FrozenLakeEnv` class. Movement boundaries are enforced so that invalid moves keep the agent in the same position.

| Component            | Design Used                              |
|----------------------|------------------------------------------|
| State representation | Single integer: state = row * 8 + column |
| Actions              | 0=Left, 1=Down, 2=Right, 3=Up            |
| Rewards              | Frozen/Start=-1, Hole=-100, Goal=+100    |
| Terminal states      | Hole and Goal states                     |

### 3. Q-Learning Algorithm

The agent uses Q-Learning, a value-based reinforcement learning algorithm. A Q-table stores the expected future reward for each state-action pair. The table is initialized to zero and updated after every transition using the Bellman Q-learning update:

$$Q(s,a) \leftarrow Q(s,a) + \alpha [r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$$

The exploration strategy is epsilon-greedy. At first the agent explores frequently, then epsilon decays over time so that the learned policy is increasingly exploited.

### 4. Training Methodology

| Hyperparameter        | Value  |
|-----------------------|--------|
| Episodes              | 20000  |
| Max steps per episode | 200    |
| Learning rate alpha   | 0.1    |
| Discount factor gamma | 0.99   |
| Initial epsilon       | 1.0    |
| Minimum epsilon       | 0.01   |
| Epsilon decay         | 0.9995 |

Training statistics recorded include episode rewards, epsilon values, successful episodes and success rate. Bonus Option B was implemented by plotting reward, moving average reward, epsilon decay and moving average success rate.

### 5. Experimental Results

| Metric                        | Result  |
|-------------------------------|---------|
| Training success rate         | 88.17%  |
| Successful training episodes  | 17634   |
| Average training reward       | 61.99   |
| Last 100 episode success rate | 98.00%  |
| Evaluation episodes           | 200     |
| Evaluation success rate       | 100.00% |
| Average evaluation reward     | 87.00   |
| Successful runs               | 200     |
| Failures                      | 0       |

### 6. Learned Policy

The learned policy below shows the recommended action for every non-terminal state. H represents a hole and G represents the goal.

```

↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓
→ → → → ↓ ↓ ↓ ↓
↑ ↑ ↑ H → → ↓ ↓
→ ↑ ↑ H → → → ↓
↑ ↑ ↑ H → → → ↓
↑ H H → → ↑ H ↓
↓ H → ↑ H ↑ H ↓
→ → ↑ H → ↑ → G




```

## 7. Hyperparameter Experiment

| Alpha | Gamma | Decay  | Eval Success | Eval Avg Reward |
|-------|-------|--------|--------------|-----------------|
| 0.10  | 0.90  | 0.9995 | 100.00%      | 87.00           |
| 0.10  | 0.95  | 0.9995 | 100.00%      | 87.00           |
| 0.10  | 0.99  | 0.9995 | 100.00%      | 87.00           |
| 0.20  | 0.99  | 0.9995 | 100.00%      | 87.00           |

## 8. Bonus Task: Training Visualizations

The following graphs were generated automatically from the training statistics. They provide visual evidence of learning progress and exploration decay.





## 9. Challenges Encountered

The main challenges were choosing stable rewards, balancing exploration and exploitation, and ensuring the agent converged to a safe route. Using infinite rewards was avoided because it can make the Q-table numerically unstable. Finite rewards allowed reliable Q-value updates.

## 10. Conclusion

The project successfully implements Frozen Lake from first principles using Q-Learning. The final greedy policy reached the goal in all 200 evaluation episodes, producing a success rate of 100.00% and zero failures. The code, README, report and results folder satisfy the assignment requirements and include the bonus visualization task.